

Cervical Cancer Classification

Loading Python Libraries

```
In [1]: import numpy as np
import pandas as pd
from typing import List
import missingno as msno
from sklearn.svm import SVC
import plotly.express as px
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from sklearn.pipeline import Pipeline
from sklearn.decomposition import PCA
from sklearn.impute import SimpleImputer
from plotly.subplots import make_subplots
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedShuffleSplit
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.metrics import precision_recall_fscore_support
from sklearn.preprocessing import RobustScaler, StandardScaler
from sklearn.ensemble import RandomForestClassifier, VotingClassifier

from imblearn.over_sampling import SMOTE, ADASYN
from imblearn.over_sampling import RandomOverSampler
```

Importing the Data

```
In [2]: # Let's fetch the youtube spam comments dataset into a data_frame
cancer_dataframe = pd.read_csv('risk_factors_cervical_cancer.csv')

# Let's take a quick look at the shape of the dataframe
print("Cancer Risk factors data shape -->", cancer_dataframe.shape)
print()

# Let's take a brief look at the contents of the dataframe
cancer_dataframe.head().style.background_gradient()
```

Cancer Risk factors data shape --> (858, 36)

Out [2]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/ year)	Hormonal Contraceptives	Con
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	
2	34	1.0	?	1.0	0.0	0.0	0.0	0.0	
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	

Gaining Insights from Data

```
In [3]: # Let's obtain a brief overview of the dataframe
cancer_dataframe.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 858 entries, 0 to 857
```

```
Data columns (total 36 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	858 non-null	int64
1	Number of sexual partners	858 non-null	object
2	First sexual intercourse	858 non-null	object
3	Num of pregnancies	858 non-null	object
4	Smokes	858 non-null	object
5	Smokes (years)	858 non-null	object
6	Smokes (packs/year)	858 non-null	object
7	Hormonal Contraceptives	858 non-null	object
8	Hormonal Contraceptives (years)	858 non-null	object
9	IUD	858 non-null	object
10	IUD (years)	858 non-null	object
11	STDs	858 non-null	object
12	STDs (number)	858 non-null	object
13	STDs:condylomatosis	858 non-null	object
14	STDs:cervical condylomatosis	858 non-null	object
15	STDs:vaginal condylomatosis	858 non-null	object
16	STDs:vulvo-perineal condylomatosis	858 non-null	object
17	STDs:syphilis	858 non-null	object
18	STDs:pelvic inflammatory disease	858 non-null	object
19	STDs:genital herpes	858 non-null	object
20	STDs:molluscum contagiosum	858 non-null	object
21	STDs:AIDS	858 non-null	object
22	STDs:HIV	858 non-null	object
23	STDs:Hepatitis B	858 non-null	object
24	STDs:HPV	858 non-null	object
25	STDs: Number of diagnosis	858 non-null	int64
26	STDs: Time since first diagnosis	858 non-null	object
27	STDs: Time since last diagnosis	858 non-null	object
28	Dx:Cancer	858 non-null	int64
29	Dx:CIN	858 non-null	int64
30	Dx:HPV	858 non-null	int64
31	Dx	858 non-null	int64
32	Hinselmann	858 non-null	int64
33	Schiller	858 non-null	int64
34	Citology	858 non-null	int64
35	Biopsy	858 non-null	int64

```
dtypes: int64(10), object(26)
```

```
memory usage: 241.4+ KB
```

```
In [4]: # Let's see descriptive statistics for all numeric columns
cancer_dataframe.describe()
```

Out [4]:

	Age	STDs: Number of diagnosis	Dx:Cancer	Dx:CIN	Dx:HPV	Dx	Hinselm
count	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000
mean	26.820513	0.087413	0.020979	0.010490	0.020979	0.027972	0.040
std	8.497948	0.302545	0.143398	0.101939	0.143398	0.164989	0.197
min	13.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
25%	20.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
50%	25.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
75%	32.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000
max	84.000000	3.000000	1.000000	1.000000	1.000000	1.000000	1.000

Observations

1. The dataset mainly covers young people with an average age of approximately 26.82 years, yet with a broad age spectrum (13 to 84 years).
2. Diagnosed conditions and positive test outcomes for STDs, cancer, HPV, and related health screenings are remarkably low, with most participants having zero diagnoses, reflecting either a generally healthy population or limited diagnostic testing among the subjects.

Data Cleaning and Processing

```
In [5]: # Let's identify columns of type object
object_cols = cancer_dataframe.select_dtypes(include=["object"]).columns
cancer_dataframe[object_cols] = cancer_dataframe[object_cols].apply(pd.to_numeric, errors='coerce')

# Let's fill NaN values with np.nan
cancer_dataframe[object_cols].fillna(np.nan, inplace=True)

# Initialize the SimpleImputer to replace missing values with the median of
imp = SimpleImputer(strategy="median")

# Let's fit the imputer and transform the data, imputing missing values
X = imp.fit_transform(cancer_dataframe)
cancer_dataframe = pd.DataFrame(X, columns=cancer_dataframe.columns)
```

<ipython-input-5-85b3a7dd27b4>:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
cancer_dataframe[object_cols].fillna(np.nan, inplace=True)

```
In [6]: def categorize_age(age):
        """
        Categorize age into predefined age groups.
        """
        if age < 12:
            category = "Child"
        elif age < 20:
            category = "Teen"
        elif age < 30:
            category = "20s"
        elif age < 40:
            category = "30s"
        elif age < 50:
            category = "40s"
        elif age < 60:
            category = "50s"
        elif age < 70:
            category = "60s"
        else:
            category = "70+"
        return category

cancer_dataframe["Age"] = cancer_dataframe["Age"].apply(int)
# Let's apply the age categorization function to each age in the dataframe
cancer_dataframe["Age_Group"] = cancer_dataframe["Age"].apply(categorize_age)
```

```
In [7]: std_columns = {
        'STDs:condylomatosis',
        'STDs:cervical condylomatosis',
        'STDs:vaginal condylomatosis',
        'STDs:vulvo-perineal condylomatosis',
        'STDs:syphilis',
        'STDs:pelvic inflammatory disease',
        'STDs:genital herpes',
        'STDs:molluscum contagiosum',
        'STDs:AIDS',
        'STDs:HIV',
        'STDs:Hepatitis B',
        'STDs:HPV'
    }

# Let's calculate the total number of STDs for each row and add it as a new
cancer_dataframe["Total_STD_Count"] = cancer_dataframe[list(std_columns)].sum(axis=1)
std_summary_by_age = cancer_dataframe.groupby("Age_Group")[list(std_columns)].sum()
```

```
In [8]: test_cols = ["Hinselmann", "Schiller", "Citology", "Biopsy"]
cancer_dataframe["total_tests"] = cancer_dataframe[test_cols].sum(axis = 1)
```

```
In [9]: # Combine general columns with STD-specific columns into a single set for co
columns_to_convert_to_int = {
    "total_tests",
    "Total_STD_Count",
    "Smokes",
    "Biopsy",
    "Dx:Cancer",
    "Num of pregnancies",
    "Number of sexual partners",
    "First sexual intercourse",
    "Hormonal Contraceptives",
    "IUD",
    "STDs",
    "STDs (number)",
    "STDs: Number of diagnosis",
    "Dx:CIN",
    "Dx:HPV",
    "Dx",
    "Hinselmann",
    "Schiller",
    "Citology"
}

columns_to_convert_to_int.update(std_columns)

for column_name in columns_to_convert_to_int:
    # Use .apply(pd.to_numeric, errors='coerce') to handle non-numeric values
    cancer_dataframe[column_name] = pd.to_numeric(cancer_dataframe[column_name], errors='coerce')
```

Check Duplicate Values

```
In [10]: # Let's check for the duplicate values
cancer_dataframe[cancer_dataframe.duplicated(keep=False)][5:]
```

```
Out[10]:
```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/ year)	Hormonal Contraceptives	C
1	15	1	14	1	0	0.0	0.0	0	
60	34	3	19	3	0	0.0	0.0	1	
66	34	3	19	3	0	0.0	0.0	1	
143	28	1	19	2	0	0.0	0.0	1	
206	28	2	20	2	0	0.0	0.0	1	

5 rows x 39 columns

```
In [11]: # Let's take a quick look at the shape of the dataframe
print("Cervical cancer data shape -->", cancer_dataframe.shape)
print()

# Let's take a brief look at the contents of the dataframe
cancer_dataframe.head().style.background_gradient()
```

Cervical cancer data shape --> (858, 39)

Out[11]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/ year)	Hormonal Contraceptives
0	18	4	15	1	0	0.000000	0.000000	0
1	15	1	14	1	0	0.000000	0.000000	0
2	34	1	17	1	0	0.000000	0.000000	0
3	52	5	16	4	1	37.000000	37.000000	1
4	46	3	21	4	0	0.000000	0.000000	1

```
In [12]: # Function to get unique_counts based on specific column
def value_counts(column_name):
    return cancer_dataframe.loc[
        :, column_name
    ].value_counts() # Returns the unique value counts
```

```
In [13]: # Iterates over all the columns of the dataframe and calls value_counts func
required_columns = [
    "total_tests",
    "Total_STD_Count",
    "Smokes",
    "Biopsy",
    "Dx:Cancer",
    "Hormonal Contraceptives",
    "IUD",
    "STDs",
    "STDs (number)",
    "STDs: Number of diagnosis",
    "Dx:CIN",
    "Dx:HPV",
    "Dx",
    "Hinselmann",
    "Schiller",
    "Citology",
]
for column_name in required_columns:
    print(f"Value Counts of {column_name}")
    value_counts_data = value_counts(column_name=column_name)
    for index, value in value_counts_data.items():
        print(f"{column_name} {index}: {value}")
    print()
```

Value Counts of total_tests

total_tests 0: 756
total_tests 1: 41
total_tests 3: 33
total_tests 2: 22
total_tests 4: 6

Value Counts of Total_STD_Count

Total_STD_Count 0: 779
Total_STD_Count 2: 37
Total_STD_Count 1: 34
Total_STD_Count 3: 7
Total_STD_Count 4: 1

Value Counts of Smokes

Smokes 0: 735
Smokes 1: 123

Value Counts of Biopsy

Biopsy 0: 803
Biopsy 1: 55

Value Counts of Dx:Cancer

Dx:Cancer 0: 840
Dx:Cancer 1: 18

Value Counts of Hormonal Contraceptives

Hormonal Contraceptives 1: 589
Hormonal Contraceptives 0: 269

Value Counts of IUD

IUD 0: 775
IUD 1: 83

Value Counts of STDs

STDs 0: 779
STDs 1: 79

Value Counts of STDs (number)

STDs (number) 0: 779
STDs (number) 2: 37
STDs (number) 1: 34
STDs (number) 3: 7
STDs (number) 4: 1

Value Counts of STDs: Number of diagnosis

STDs: Number of diagnosis 0: 787
STDs: Number of diagnosis 1: 68
STDs: Number of diagnosis 2: 2
STDs: Number of diagnosis 3: 1

Value Counts of Dx:CIN

Dx:CIN 0: 849
Dx:CIN 1: 9

Value Counts of Dx:HPV

Dx:HPV 0: 840
Dx:HPV 1: 18

Value Counts of Dx

Dx 0: 834

Dx 1: 24

Value Counts of Hinselmann

Hinselmann 0: 823

Hinselmann 1: 35

Value Counts of Schiller

Schiller 0: 784

Schiller 1: 74

Value Counts of Citology

Citology 0: 814


Citology 1: 44

See how many missing data points we have

```
In [14]: cancer_dataframe.isnull().sum()
```

```
Out[14]: Age                                0
Number of sexual partners                  0
First sexual intercourse                    0
Num of pregnancies                         0
Smokes                                     0
Smokes (years)                            0
Smokes (packs/year)                       0
Hormonal Contraceptives                   0
Hormonal Contraceptives (years)           0
IUD                                         0
IUD (years)                               0
STDs                                       0
STDs (number)                             0
STDs:condylomatosis                       0
STDs:cervical condylomatosis              0
STDs:vaginal condylomatosis              0
STDs:vulvo-perineal condylomatosis        0
STDs:syphilis                             0
STDs:pelvic inflammatory disease          0
STDs:genital herpes                       0
STDs:molluscum contagiosum                0
STDs:AIDS                                 0
STDs:HIV                                  0
STDs:Hepatitis B                          0
STDs:HPV                                  0
STDs: Number of diagnosis                 0
STDs: Time since first diagnosis           0
STDs: Time since last diagnosis           0
Dx:Cancer                                 0
Dx:CIN                                    0
Dx:HPV                                    0
Dx                                         0
Hinselmann                                0
Schiller                                  0
Citology                                  0
Biopsy                                    0
Age_Group                                 0
Total_STD_Count                           0
total_tests                               0
dtype: int64
```

```
In [15]: msno.bar(cancer_dataframe, figsize = (16,5),color = "pink")
plt.show()
```



Data Visualization

```
In [16]: # Let's visualize the distribution of age
age_dist = px.histogram(cancer_dataframe, x="Age", marginal="box", color_discrete_sequence=px.colors.qualitative.M10)
age_dist.update_layout(title="Age distribution")
age_dist.show()
```

```
In [17]: diagnoses_num_partner_compare_cols = ['Dx:Cancer',  
                                                'Dx:HPV',  
                                                "Number of sexual partners",]  
corr_matrix = cancer_dataframe[diagnoses_num_partner_compare_cols].corr()  
diagnoses_num_partner_heatmap = px.imshow(corr_matrix,  
                                           aspect="auto",  
                                           color_continuous_scale="gnbu",  
                                           text_auto=True)  
diagnoses_num_partner_heatmap.show()
```

Imbalance Data

```
In [18]: dx_cancer = px.histogram(cancer_dataframe, y="Dx:Cancer")
dx_cancer.update_layout(bargap=0.2)
dx_cancer.update_layout(title = "Imbalanced Classes")
dx_cancer.show()
```

```
In [19]: X = cancer_dataframe.drop(["Dx:Cancer", "Age_Group"], axis=1)
y = cancer_dataframe["Dx:Cancer"].copy()
```

```
In [20]: # Let's initialize ADASYN sample
adaptive_sampling = ADASYN(random_state=42)

# Let's apply ADASYN to resample the features (X) and target (y)
features_resampled, target_resampled = adaptive_sampling.fit_resample(X, y)

cancer_dataframe = features_resampled.join(target_resampled, how='outer')
```

```
In [21]: # Let's apply the age categorization function to each age in the dataframe
cancer_dataframe["Age_Group"] = cancer_dataframe["Age"].apply(categorize_age)
```

```
In [22]: # Let's create a histogram to visualize the distribution of Dx:Cancer in the
cancer_diagnosis_distribution = px.histogram(cancer_dataframe, y="Dx:Cancer")
cancer_diagnosis_distribution.update_layout(
    bargap=0.2, # Adjust the gap between bars
    title_text="Distribution of Cancer Diagnosis - Balanced Classes" # Update title
)

# Let's display the histogram
cancer_diagnosis_distribution.show()
```

```
In [23]: def prepare_data(df, target, stratify_col, test_ratio=0.2, random_seed=62, c
        split = StratifiedShuffleSplit(n_splits=1, test_size=test_ratio, random_
        for train_idx, test_idx in split.split(df, df[stratify_col]):
            train_set = df.loc[train_idx]
            test_set = df.loc[test_idx]

            if cols_to_remove:
                for set_ in (train_set, test_set):
                    set_.drop(cols_to_remove, axis=1, inplace=True)

            X_train = train_set.drop(target, axis=1)
            y_train = train_set[target].copy()

            X_test = test_set.drop(target, axis=1)
            y_test = test_set[target].copy()

            return X_train, y_train, X_test, y_test

cols_to_drop = ["Age_Group", "Total_STD_Count", "total_tests"]
X_train, y_train, X_test, y_test = prepare_data(cancer_dataframe, "Dx:Cancer
```

PCA and Model

```
In [24]: def preprocess_features(feature_data, n_pca_components=13, pipeline=None):
    if pipeline is None:
        # Define and fit a new preprocessing pipeline if none is provided
        pipeline = Pipeline([
            ("data_scaler", RobustScaler()),
            ("dimensionality_reducer", PCA(n_components=n_pca_components))
        ])
        transformed_features = pipeline.fit_transform(feature_data)
    else:
        # Use the provided pipeline to transform the data
        transformed_features = pipeline.transform(feature_data)

    return transformed_features, pipeline

X_train_preprocessed, training_pipeline = preprocess_features(X_train, n_pca
X_test_preprocessed, _ = preprocess_features(X_test, pipeline=training_pipeline)
```

```
In [25]: from sklearn.pipeline import make_pipeline

# Let's define a parameter grid for hyperparameter tuning
parameter_grid = {'logisticregression__C': np.logspace(-5, 8, 15)}

# Let's create a pipeline that first scales the data
pipeline = make_pipeline(StandardScaler(), LogisticRegression(max_iter=1000))

# Set up GridSearchCV with the pipeline
grid_search = GridSearchCV(pipeline, parameter_grid, cv=10, refit=True)

# Let's fit the grid search to the training data
logistic_regression_cv = grid_search.fit(X_train_preprocessed, y_train)
```

Random Forest Classifier

```
In [26]: random_forest_classifier = RandomForestClassifier(n_estimators=2500, max_lea
max_features="sqrt")
```

KNN Classifier


```
In [27]: # Let's define KNNClassifier

knn_classifier = KNeighborsClassifier()
knn_param_grid = {"n_neighbors": list(np.arange(1, 100, 2))}
knn_classifier_cv = GridSearchCV(knn_classifier, knn_param_grid, cv=10, refit=True)
knn_classifier_cv = KNeighborsClassifier(**knn_classifier_cv.best_params_)
```

SVM Classifier

```
In [28]: # Let's define SVM Classifier

svm_classifier = SVC()
svc_param_grid = {"C": np.logspace(-3, 2, 6), "gamma": np.logspace(-3, 2, 6)}
svm_classifier_cv = GridSearchCV(svm_classifier, svc_param_grid, cv=5)
```

```
In [29]: col_names = ["Classifier Name", "Accuracy Score", "Precision Score",
                      "Recall Score", "F1 Score"]
summary_df = pd.DataFrame(columns=col_names)

est_name = []
est_acc = []
precision_score = []
recall_score = []
f1score = []
est_conf_matrix = []

estimators = [
    ("LogisticRegression", logistic_regression_cv),
    ("RandomForestClassifier", random_forest_classifier),
    ("KNeighborsClassifier", knn_classifier_cv),
    ("SupportVectorClassifier", svm_classifier_cv)]

for i in range(0, len(estimators)):
    clf_name = estimators[i][0]
    clf = estimators[i][1]
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    # print(pd.crosstab(y_test, y_pred, rownames=["Actual"], colnames=["predicted"]))
    est_name.append(estimators[i][0])
    est_acc.append(accuracy_score(y_test, y_pred))
    scores = precision_recall_fscore_support(y_test, y_pred, average="weighted")
    precision_score.append(scores[0])
    recall_score.append(scores[1])
    f1score.append(scores[2])
    est_conf_matrix.append(confusion_matrix(y_test, y_pred))

summary_df[col_names[0]] = est_name
summary_df[col_names[1]] = est_acc
summary_df[col_names[2]] = precision_score
summary_df[col_names[3]] = recall_score
summary_df[col_names[4]] = f1score
```

Results

```
In [30]: # Let's print the results of LogisticRegression, Random Forest, KNN and SVC
summary_df
```

```
Out[30]:
```

	Classifier Name	Accuracy Score	Precision Score	Recall Score	F1 Score
0	LogisticRegression	1.000000	1.000000	1.000000	1.000000
1	RandomForestClassifier	0.991071	0.991226	0.991071	0.991070
2	KNeighborsClassifier	0.955357	0.959057	0.955357	0.955292
3	SupportVectorClassifier	1.000000	1.000000	1.000000	1.000000

```
In [31]: # Define a list of color scales for visual variety
unique_color_scales = ["agsunset", "teal", "purples", "viridis"]

# Iterate over the confusion matrices and their corresponding names
for index, conf_matrix in enumerate(est_conf_matrix):
    # Generate a heatmap for each confusion matrix
    heatmap_visual = px.imshow(conf_matrix, aspect="auto", text_auto=True,
                               color_continuous_scale=unique_color_scales[index])

    # Customize the layout with dynamic titles and axis labels
    heatmap_visual.update_layout(title=est_name[index])
    heatmap_visual.update_xaxes(title="Predicted Labels")
    heatmap_visual.update_yaxes(title="Actual Labels")

    # Display the heatmap
    heatmap_visual.show()
```